

Setup

- install Java; check version (`java -version`)
- download Confluent platform from <https://www.confluent.io/download/>
- unzip and untar (`tar xvfz confluent-5*.tar.gz`)
- move to /opt
- create symbolic link (`ln -s confluent-5* confluent`)
- add to path (`export PATH=${PATH}:/opt/confluent/bin`)
- install the Confluent CLI (see <https://docs.confluent.io/current/cli/installing.html#scripted-installation>)
 - `curl -L https://cnfl.io/cli | sh -s -- -b /opt/confluent/bin`
- Start
 - `confluent local start ksql-server`
 - or if using Docker `docker-compose up -d`

Confluent CLI changes from version 5.3

The *Confluent CLI* changed significantly in version 5.3. The most important change is the inclusion of the `local` paramter when interacting with a local development environment.

For example, to start the ksql-server

- Prior to 5.3 : `confluent start ksql-server`
- From 5.3 : `confluent local start ksql-server`

Create a topic

```
kafka-topics --zookeeper localhost:2181 --create --partitions 1 --
replication-factor 1 --topic USERS
```

Or you can use the more modern syntax

```
kafka-topics --bootstrap-server localhost:9092 --create --partitions 1 --
replication-factor 1 --topic USERS
```

Get started - KSQL Command Line

At UNIX prompt

```
kafka-console-producer --broker-list localhost:9092 --topic USERS << EOF
Alice,US
```

```
Bob,GB
Carol,AU
Dan,US
EOF
```

At KSQL prompt

```
show topics;

-- this will show nothing
print 'USERS';

print 'USERS' from beginning;

print 'USERS' from beginning limit 2;

print 'USERS' from beginning interval 2 limit 2 ;
```

Get started - Create a stream with CSV

Special note for 5.4 onwards (ksqlDB)

There are two categories of queries :

- **Push queries:** query the state of the system in motion and continue to output results until they meet a LIMIT condition or are terminated by the user. This was the default behavior in older versions of KSQL. 'EMIT CHANGES' is used to indicate a query is a push query.
- **Pull queries:** query the current state of the system, return a result, and terminate. Use this to select a result as of "now". New from 5.4. KSQL currently only supports pull queries on materialized aggregate tables (sometimes referred to as *materialized views*). i.e. those created by a 'CREATE TABLE AS SELECT', <aggregate_functions> FROM GROUP BY ' style statement. A query must use a predicate against **ROWKEY**

At KSQL prompt

```
create stream users_stream (name VARCHAR, countrycode VARCHAR) WITH
(KAFKA_TOPIC='USERS', VALUE_FORMAT='DELIMITED');

list streams;

-- nothing will get shown
-- 5.3 and earlier
select name, countrycode from users_stream;

-- 5.4 onwards
select name, countrycode from users_stream emit changes;
```

auto.offset.reset - Determines what to do when there is no initial offset in Apache Kafka or if the current offset does not exist on the server. The default value in KSQL is *latest*, which means all Kafka topics are read from the latest available offset. For example, to change it to *earliest* by using the KSQL command line:

```
-- default to beginning of time
SET 'auto.offset.reset'='earliest';

-- now will see something
-- 5.3 and earlier
select name, countrycode  from users_stream;

-- 5.4 onwards
select name, countrycode  from users_stream emit changes;

-- stop after 4 records
-- 5.3 and earlier
select name, countrycode  from users_stream limit 4;

-- 5.4 onwards
select name, countrycode  from users_stream emit changes limit 4;

-- basic aggregate
-- 5.3 and earlier
select countrycode, count(*) from users_stream group by countrycode;

-- 5.4 onwards
select countrycode, count(*) from users_stream group by countrycode emit
changes;

drop stream if exists users_stream delete topic;

list streams;

show topics;
```

Create a stream with JSON

At UNIX prompt

```
kafka-topics --zookeeper localhost:2181 --create --partitions 1 --
replication-factor 1 --topic USERPROFILE

kafka-console-producer --broker-list localhost:9092 --topic USERPROFILE <<
EOF
{"userid": 1000, "firstname":"Alison", "lastname":"Smith",
```

```
"countrycode":"GB", "rating":4.7}
EOF

kafka-console-producer --broker-list localhost:9092 --topic USERPROFILE <<
EOF
{"userid": 1001, "firstname":"Bob", "lastname":"Smith",
"countrycode":"US", "rating":4.2}
EOF
```

At KSQL prompt

```
CREATE STREAM userprofile (userid INT, firstname VARCHAR, lastname
VARCHAR, countrycode VARCHAR, rating DOUBLE) \
  WITH (VALUE_FORMAT = 'JSON', KAFKA_TOPIC = 'USERPROFILE');
```

```
SET 'auto.offset.reset'='earliest';

-- 5.3 and earlier
select firstname, lastname, countrycode, rating from userprofile;

-- 5.4 onwards
select firstname, lastname, countrycode, rating from userprofile emit
changes;

Alison | Smith | GB | 4.7
```

Manipulate a stream

Run a data gen

At UNIX prompt

```
ksql-datagen schema=./datagen/userprofile.avro format=json
topic=USERPROFILE key=userid maxInterval=5000 iterations=100
```

At KSQL prompt

```
-- Review a stream - every 5th row
print 'USERPROFILE' interval 5;
```

Manipulate a stream

At KSQL prompt

```
ksql> describe userprofile;
```

```
Name          : USERPROFILE
Field         | Type
-----
ROWTIME       | BIGINT          (system)  <-- NOTE
ROWKEY        | VARCHAR(STRING) (system)  <-- NOTE
USERID        | INTEGER
FIRSTNAME     | VARCHAR(STRING)
LASTNAME      | VARCHAR(STRING)
COUNTRYCODE   | VARCHAR(STRING)
RATING        | DOUBLE
```

```
-- 5.3 and earlier
```

```
select rowtime, firstname from userprofile;
```

```
-- 5.4 onwards
```

```
select rowtime, firstname from userprofile emit changes;
```

- Review **Scalar functions** at <https://docs.confluent.io/current/ksql/docs/developer-guide/syntax-reference.html#scalar-functions>

```
SELECT TIMESTAMPTOSTRING(rowtime, 'dd/MMM HH:mm') as createtime, firstname
from userprofile emit changes;
```

```
select TIMESTAMPTOSTRING(rowtime, 'dd/MMM HH:mm') as createtime,
firstname + ' ' + ucase(lastname) as full_name
from userprofile emit changes;
```

Create a stream from a stream

At KSQL prompt

```
select firstname + ' '
+ ucase( lastname)
+ ' from ' + countrycode
+ ' has a rating of ' + cast(rating as varchar) + ' stars. '
+ case when rating < 2.5 then 'Poor'
      when rating between 2.5 and 4.2 then 'Good'
      else 'Excellent'
end as description
from userprofile emit changes;
```

```
Bob FAWCETT from IN has a rating of 4.4 stars. | Excellent
Heidi COEN from US has a rating of 4.9 stars. | Excellent
Bob FAWCETT from IN has a rating of 2.2 stars. | Poor
```

At KSQL prompt

Review the script `user_profile_pretty.ksql`

```
list streams;

run script 'user_profile_pretty.ksql';

list streams;

describe extended user_profile_pretty;

select description from user_profile_pretty emit changes;

drop stream user_profile_pretty;

terminate query CSAS_USER_PROFILE_PRETTY_4;

drop stream user_profile_pretty;

list streams;

drop stream IF EXISTS user_profile_pretty DELETE TOPIC;
```

Create a table

At UNIX prompt

```
kafka-topics --zookeeper localhost:2181 --create --partitions 1 --
replication-factor 1 --topic COUNTRY-CSV

kafka-console-producer --broker-list localhost:9092 --topic COUNTRY-CSV --
property "parse.key=true" --property "key.separator=:" << EOF
AU:AU,Australia
IN:IN,India
GB:GB,UK
US:US,United States
EOF
```

At KSQL prompt

```
CREATE TABLE COUNTRYTABLE (countrycode VARCHAR, countryname VARCHAR) WITH
(KAFKA_TOPIC='COUNTRY-CSV', VALUE_FORMAT='DELIMITED', KEY =
'countrycode');

show tables;

describe COUNTRYTABLE;

describe extended COUNTRYTABLE;

SET 'auto.offset.reset'='earliest';

select countrycode, countryname from countrytable emit changes;

-- Note the countryname is "UK"
select countrycode, countryname from countrytable where countrycode='GB'
emit changes limit 1;

-- This does not exist
select countrycode, countryname from countrytable where countrycode='FR'
emit changes;
```

Update a table

One record updated (UK->United Kingdom), one record added (FR)

At UNIX prompt

```
kafka-console-producer --broker-list localhost:9092 --topic COUNTRY-CSV --
property "parse.key=true" --property "key.separator=:" << EOF
GB:GB,United Kingdom
FR:FR,France
EOF
```

At KSQL prompt

```
select countrycode, countryname from countrytable emit changes;

-- Note the countryname has changed to "United Kingdom"
select countrycode, countryname from countrytable where countrycode='GB'
emit changes limit 1;

-- And now appears
select countrycode, countryname from countrytable where countrycode='FR'
emit changes;
```

Join

Join user stream to country table

At KSQL prompt

```
select up.firstname, up.lastname, up.countrycode, ct.countryname
from USERPROFILE up
left join COUNTRYTABLE ct on ct.countrycode=up.countrycode emit changes;

create stream up_joined as
select up.firstname
+ ' ' + ucase(up.lastname)
+ ' from ' + ct.countryname
+ ' has a rating of ' + cast(rating as varchar) + ' stars.' as description
from USERPROFILE up
left join COUNTRYTABLE ct on ct.countrycode=up.countrycode;

select * from up_joined emit changes;
```

Pull Queries

Pull Query - new in ksqlDB (5.4 onwards)

```
SET 'auto.offset.reset'='earliest';

CREATE STREAM driverLocations (driverId INTEGER, countrycode VARCHAR, city
VARCHAR, driverName VARCHAR)
  WITH (kafka_topic='driverlocations', key='driverId',
value_format='json', partitions=1);

INSERT INTO driverLocations (driverId, countrycode, city, driverName)
VALUES (1, 'AU', 'Sydney', 'Alice');
INSERT INTO driverLocations (driverId, countrycode, city, driverName)
VALUES (2, 'AU', 'Melbourne', 'Bob');
INSERT INTO driverLocations (driverId, countrycode, city, driverName)
VALUES (3, 'GB', 'London', 'Carole');
INSERT INTO driverLocations (driverId, countrycode, city, driverName)
VALUES (4, 'US', 'New York', 'Derek');
```

```
create table countryDrivers as select countrycode, count(*) as numDrivers
from driverLocations group by countrycode;
```



```
-- note: as a pull query we don't use "emit"
select countrycode, numdrivers from countryDrivers where rowkey='AU';

INSERT INTO driverLocations (driverId, countrycode, city, driverName)
VALUES (5, 'AU', 'Sydney', 'Emma');
select countrycode, numdrivers from countryDrivers where rowkey='AU';

-- note: as a pull query we don't use "emit"
select countrycode, numdrivers from countryDrivers where rowkey='AU';
```

Kafka Connect with ksqldb

To run Postgres and Confluent platform

```
docker-compose up -d
```

Start ksqldb KSQL CLI

```
docker-compose exec ksqldb-cli ksql http://ksqldb-server:8088
```

Kafka Connect

```
cat postgres-setup.sql

docker-compose exec postgres psql -U postgres -f /postgres-setup.sql
```

To look at the Postgres table

```
docker-compose exec postgres psql -U postgres -c "select * from carusers;"
```

```
CREATE SOURCE CONNECTOR `postgres-jdbc-source` WITH(
  "connector.class"='io.confluent.connect.jdbc.JdbcSourceConnector',
  "connection.url"='jdbc:postgresql://postgres:5432/postgres',
  "mode"='incrementing',
  "incrementing.column.name"='ref',
  "table.whitelist"='carusers',
  "connection.password"='password',
  "connection.user"='postgres',
  "topic.prefix"='db-',
  "key"='username');
```

```
print 'db-carusers' from beginning;
```

In another window, insert a new database row

```
docker exec -it postgres psql -U postgres -c "INSERT INTO carusers  
(username) VALUES ('Derek');"
```

Data Formats

Imagine a *complaints* stream of unhappy customers. Explore the different data formats (CSV, JSON, AVRO)

Column	AVRO Type	KSQL Type
customer_name	string	VARCHAR
complaint_type	string	VARCHAR
trip_cost	float	DOUBLE
new_customer	boolean	BOOLEAN

CSV Delimited

At UNIX prompt

```
kafka-topics --zookeeper localhost:2181 --create --partitions 1 --  
replication-factor 1 --topic COMPLAINTS_CSV  
  
kafka-console-producer --broker-list localhost:9092 --topic COMPLAINTS_CSV  
<< EOF  
Alice, Late arrival, 43.10, true  
EOF
```

At KSQL prompt

```
CREATE STREAM complaints_csv (customer_name VARCHAR, complaint_type  
VARCHAR, trip_cost DOUBLE, new_customer BOOLEAN) \  
  WITH (VALUE_FORMAT = 'DELIMITED', KAFKA_TOPIC = 'COMPLAINTS_CSV');  
  
select * from complaints_csv emit changes;
```

CSV - experience with bad data

At UNIX prompt

```
kafka-console-producer --broker-list localhost:9092 --topic COMPLAINTS_CSV << EOF
Alice, Bob and Carole, Bad driver, 43.10, true
EOF
```

JSON

At UNIX prompt

```
kafka-topics --zookeeper localhost:2181 --create --partitions 1 --replication-factor 1 --topic COMPLAINTS_JSON

kafka-console-producer --broker-list localhost:9092 --topic COMPLAINTS_JSON << EOF
{"customer_name":"Alice, Bob and Carole", "complaint_type":"Bad driver", "trip_cost": 22.40, "new_customer": true}
EOF
```

At KSQL prompt

```
CREATE STREAM complaints_json (customer_name VARCHAR, complaint_type VARCHAR, trip_cost DOUBLE, new_customer BOOLEAN) \
  WITH (VALUE_FORMAT = 'JSON', KAFKA_TOPIC = 'COMPLAINTS_JSON');

select * from complaints_json;
```

JSON - experience with bad data

At UNIX prompt

```
kafka-console-producer --broker-list localhost:9092 --topic COMPLAINTS_JSON << EOF
{"customer_name":"Bad Data", "complaint_type":"Bad driver", "trip_cost": 22.40, "new_customer": ShouldBeABoolean}
EOF
```

Review the KSQL Server logs `confluent local log ksql-server`

Now look at the *KSQL Server log*. We can see bad data is noticed; but hidden in a conversion error message

```
at [Source: (byte[])"{"customer_name":"Bad Data", "complaint_type":"Bad driver", "trip_cost": 22.40, "new_customer": ShouldBeABoolean}"; line: 1, column: 105]
Caused by: com.fasterxml.jackson.core.JsonParseException: Unrecognized token 'ShouldBeABoolean': was expecting ('true', 'false' or 'null')
```

AVRO

At UNIX prompt

```
kafka-topics --zookeeper localhost:2181 --create --partitions 1 --replication-factor 1 --topic COMPLAINTS_AVRO

kafka-avro-console-producer --broker-list localhost:9092 --topic COMPLAINTS_AVRO \
--property value.schema='
{
  "type": "record",
  "name": "myrecord",
  "fields": [
    {"name": "customer_name", "type": "string" }
    , {"name": "complaint_type", "type": "string" }
    , {"name": "trip_cost", "type": "float" }
    , {"name": "new_customer", "type": "boolean"}
  ]
}' << EOF
{"customer_name":"Carol", "complaint_type":"Late arrival", "trip_cost":19.60, "new_customer": false}
EOF
```

At KSQL prompt

```
-- Note no columns or data type specified
create stream complaints_avro with (kafka_topic='COMPLAINTS_AVRO',
value_format='AVRO');

describe extended complaints_avro;
```

AVRO - experience with bad data

At UNIX prompt - note bad data is noted at serialization time

```
kafka-avro-console-producer --broker-list localhost:9092 --topic COMPLAINTS_AVRO \
--property value.schema='
{
```

```

    "type": "record",
    "name": "myrecord",
    "fields": [
      {"name": "customer_name", "type": "string" },
      , {"name": "complaint_type", "type": "string" }
      , {"name": "trip_cost", "type": "float" }
      , {"name": "new_customer", "type": "boolean"}
    ]
  }' << EOF
{"customer_name":"Bad Data", "complaint_type":"Bad driver", "trip_cost":
22.40, "new_customer": ShouldBeABoolean}
EOF

```

AVRO Schema Evolution

At UNIX prompt

```

# Optional : strart Confluent Control Center
confluent local start

curl -s -X GET http://localhost:8081/subjects/COMPLAINTS_AVRO-
value/versions

kafka-avro-console-producer --broker-list localhost:9092 --topic
COMPLAINTS_AVRO \
--property value.schema='
{
  "type": "record",
  "name": "myrecord",
  "fields": [
    {"name": "customer_name", "type": "string" }
    , {"name": "complaint_type", "type": "string" }
    , {"name": "trip_cost", "type": "float" }
    , {"name": "new_customer", "type": "boolean"}
    , {"name": "number_of_rides", "type": "int", "default" : 1}
  ]
}' << EOF
{"customer_name":"Ed", "complaint_type":"Dirty car", "trip_cost": 29.10,
"new_customer": false, "number_of_rides": 22}
EOF

curl -s -X GET http://localhost:8081/subjects/COMPLAINTS_AVRO-
value/versions

curl -s -X GET http://localhost:8081/subjects/COMPLAINTS_AVRO-
value/versions/1 | jq '.'

curl -s -X GET http://localhost:8081/subjects/COMPLAINTS_AVRO-
value/versions/2 | jq '.'

```

At KSQL prompt

```
ksql> describe complaints_avro;
```

Name : COMPLAINTS_AVRO

Field | Type

ROWTIME	BIGINT	(system)
ROWKEY	VARCHAR(STRING)	(system)
CUSTOMER_NAME	VARCHAR(STRING)	
COMPLAINT_TYPE	VARCHAR(STRING)	
TRIP_COST	DOUBLE	
NEW_CUSTOMER	BOOLEAN	

```
ksql> create stream complaints_avro_v2 with
(kafka_topic='COMPLAINTS_AVRO', value_format='AVRO');
```

```
ksql> describe complaints_avro_v2;
```

Name : COMPLAINTS_AVRO_V2

Field | Type

ROWTIME	BIGINT	(system)
ROWKEY	VARCHAR(STRING)	(system)
CUSTOMER_NAME	VARCHAR(STRING)	
COMPLAINT_TYPE	VARCHAR(STRING)	
TRIP_COST	DOUBLE	
NEW_CUSTOMER	BOOLEAN	
NUMBER_OF_RIDES	INTEGER	

<-- *** NOTE new column

Nested JSON

Imagine we have data like this

```
{
  "city": {
    "name": "Sydney",
    "country": "AU",
    "latitude": -33.8688,
    "longitude": 151.2093
  },
  "description": "light rain",
  "clouds": 92,
  "deg": 26,
  "humidity": 94,
```

```
"pressure": 1025.12,  
"rain": 1.25  
}
```

At UNIX prompt

```
kafka-topics --zookeeper localhost:2181 --create --partitions 1 --  
replication-factor 1 --topic WEATHERNESTED  
  
cat demo-weather.json | kafka-console-producer --broker-list  
localhost:9092 --topic WEATHERNESTED
```

Extract like this - At KSQL prompt

```
SET 'auto.offset.reset'='earliest';  
  
CREATE STREAM weather  
  (city STRUCT <name VARCHAR, country VARCHAR, latitude DOUBLE,  
  longitude DOUBLE>,  
   description VARCHAR,  
   clouds BIGINT,  
   deg BIGINT,  
   humidity BIGINT,  
   pressure DOUBLE,  
   rain DOUBLE)  
WITH (KAFKA_TOPIC='WEATHERNESTED', VALUE_FORMAT='JSON');  
  
SELECT city->name AS city_name, city->country AS city_country, city-  
>latitude as latitude, city->longitude as longitude, description, rain  
from weather emit changes;
```

Build rekeyed table

- create a table based on rekeyed **city** field from **weather** stream
- At KSQL prompt

```
create stream weatherraw with (value_format='AVRO') as SELECT city->name  
AS city_name, city->country AS city_country, city->latitude as latitude,  
city->longitude as longitude, description, rain from weather ;  
  
list streams;  
-- note AVRO  
  
describe extended weatherraw;
```

Now notice the *Key field*

```
ksql> describe extended weatherraw;
>

Name           : WEATHERRAW
Type           : STREAM
Key field      :                <- *** NOTE BLANK ***
Key format     : STRING
Timestamp field : Not set - using <ROWTIME>
Value format   : AVRO
Kafka topic    : WEATHERRAW (partitions: 4, replication: 1)
```

```
create stream weatherrekeyed as select * from weatherraw partition by
city_name;

describe extended weatherrekeyed;
```

Now notice the *Key field*

```
ksql> describe extended weatherrekeyed;
>

Name           : WEATHERREKEYED
Type           : STREAM
Key field      : CITY_NAME      <- *** Keyed on city ***
Key format     : STRING
Timestamp field : Not set - using <ROWTIME>
Value format   : AVRO
Kafka topic    : WEATHERREKEYED (partitions: 4, replication: 1)
```

```
create table weathernow with (kafka_topic='WEATHERREKEYED',
value_format='AVRO', key='CITY_NAME');

select * from weathernow emit changes;

select * from weathernow where city_name = 'San Diego' emit changes;
```

Let's make it sunny! At UNIX prompt

```
cat demo-weather-changes.json | kafka-console-producer --broker-list
localhost:9092 --topic WEATHERNESTED
```


At KSQL prompt

```
select * from weathernow where city_name = 'San Diego' emit changes;
```

Repartition

When you use KSQL to join streaming data, you must ensure that your streams and tables are co-partitioned, which means that input records on both sides of the join have the same configuration settings for partitions.

At UNIX prompt

```
kafka-topics --zookeeper localhost:2181 --create --partitions 2 --
replication-factor 1 --topic DRIVER_PROFILE

kafka-console-producer --broker-list localhost:9092 --topic DRIVER_PROFILE
<< EOF
{"driver_name":"Mr. Speedy", "countrycode":"AU", "rating":2.4}
EOF
```

At KSQL prompt

```
CREATE STREAM DRIVER_PROFILE (driver_name VARCHAR, countrycode VARCHAR,
rating DOUBLE)
  WITH (VALUE_FORMAT = 'JSON', KAFKA_TOPIC = 'DRIVER_PROFILE');
```

```
select dp.driver_name, ct.countryname, dp.rating
from DRIVER_PROFILE dp
left join COUNTRYTABLE ct on ct.countrycode=dp.countrycode emit changes;
```

Can't join DRIVER_PROFILE with COUNTRYTABLE since the number of partitions don't match. DRIVER_PROFILE partitions = 2; COUNTRYTABLE partitions = 1. Please repartition either one so that the number of partitions match.

We can fix this by co-partitioning, use the PARTITION BY clause. At KSQL prompt

```
create stream driverprofile_rekeyed with (partitions=1) as select * from
DRIVER_PROFILE partition by driver_name;

select dp2.driver_name, ct.countryname, dp2.rating
```

```
from DRIVERPROFILE_REKEYED dp2
left join COUNTRYTABLE ct on ct.countrycode=dp2.countrycode emit changes;
```

Mergin Streams; Concat Topics with INSERT

- create stream of requested rides in Europe using data gen
- create stream of requested rides in USA using data gen
- combine into single stream of all requested rides using *INSERT*

At UNIX prompt

```
ksql-datagen schema=./datagen/riderequest-europe.avro format=avro
topic=riderequest-europe key=rideid maxInterval=5000 iterations=100

ksql-datagen schema=./datagen/riderequest-america.avro format=avro
topic=riderequest-america key=rideid maxInterval=5000 iterations=100
```

At KSQL prompt

```
create stream rr_america_raw with (kafka_topic='riderequest-america',
value_format='avro');

create stream rr_europe_raw with (kafka_topic='riderequest-europe',
value_format='avro');

select * from rr_america_raw emit changes;

select * from rr_europe_raw emit changes;

create stream rr_world as select 'Europe' as data_source, * from
rr_europe_raw;

insert into rr_world      select 'Americas' as data_source, * from
rr_america_raw;

select * from rr_world emit changes;
```

Windows

- how many requests are arriving each time period
- At KSQL prompt

```
select data_source, city_name, count(*)
from rr_world
window tumbling (size 60 seconds)
group by data_source, city_name emit changes;
```

```
select data_source, city_name, COLLECT_LIST(user)
from rr_world
window tumbling (size 60 seconds)
group by data_source, city_name emit changes;
```

```
select data_source, city_name, COLLECT_LIST(user)
from rr_world WINDOW SESSION (60 SECONDS)
group by data_source, city_name emit changes;

select TIMESTAMPTOSTRING(WindowStart(), 'HH:mm:ss')
, TIMESTAMPTOSTRING(WindowEnd(), 'HH:mm:ss')
, data_source
, TOPK(city_name, 3)
, count(*)
FROM rr_world
WINDOW TUMBLING (SIZE 1 minute)
group by data_source
emit changes;
```

Geospacial

- create stream - distance of car to waiting rider
- At KSQL prompt

```
select * from rr_world emit changes;

describe rr_world;

create stream requested_journey as
select rr.latitude as from_latitude
, rr.longitude as from_longitude
, rr.user
, rr.city_name as city_name
, w.city_country
, w.latitude as to_latitude
, w.longitude as to_longitude
, w.description as weather_description
, w.rain
from rr_world rr
```

```
left join weathernow w on rr.city_name = w.city_name;

create stream ridetodest as
select user
, city_name
, city_country
, weather_description
, rain
, GEO_DISTANCE(from_latitude, from_longitude, to_latitude, to_longitude,
'km') as dist
from requested_journey;
```

```
select user + ' is travelling ' + cast(round(dist) as varchar) + ' km to '
+ city_name + ' where the weather is reported as ' + weather_description
from ridetodest emit changes;
```

Alice is at (52,0) and is travelling 215 km to Manchester where it is
SUNNY

Heidi is at (51,-1) and is travelling 88 km to London where it is heavy
rain

Grace is at (50,-1) and is travelling 138 km to London where it is heavy
rain

UDF - Build and deploy KSQL User Defined Anomaly Functions

- write a UDF to calculate drive time based on
 - distance to travel
 - weather conditions

Compile Code to Create Anomaly Functions

- Have a look at the file `java/src/main/java/com/vsimon/kafka/streams/TaxiWait.java`
- If you don't want to compile the code; just copy the JAR from `java/pre-compiled/ksql-udf-taxi-1.0.jar`
- Download Maven and follow the installation instructions (<https://maven.apache.org/>)

```
cd java
mvn clean package
ls target/ksql-udf-taxi-1.0.jar
```

Deploy KSQL User Defined Functions

Find the location of your extension directory. From KSQL

```
ksql> LIST PROPERTIES;
```

Property	Effective Value

ksql.extension.dir	ext <-- *** Look for this

```
# Stop (just the) KSQL-Server
confluent local stop ksql-server

# Create an ext (extensions) directory in ${CONFLUENT_HOME}/ext
mkdir /opt/confluent/ext

# build ksql-udf-taxi.jar as above and copy into ext directory
cp target/ksql-udf-taxi-1.0.jar /opt/confluent/ext

# or to use the pre-compile one
cp pre-compiled/ksql-udf-taxi-1.0.jar /opt/confluent/ext

# Restart KSQL server
confluent local start ksql-server
```

Check KSQL User Defined Functions Available

Start **ksql** client and verify

```
ksql> list functions;
```

Function Name	Type

SUM	AGGREGATE
TAXI_WAIT	SCALAR <--- You need this one
TIMESTAMP TO STRING	SCALAR

```
ksql> DESCRIBE FUNCTION TAXI_WAIT;
```

```
Name      : TAXI_WAIT
Overview  : Return expected wait time in minutes
Type      : scalar
Jar       : /etc/ksql/ext/ksql-udf-taxi-1.0.jar
Variations :

          Variation : TAXI_WAIT(VARCHAR, DOUBLE)
```

```
Returns      : DOUBLE
Description  : Given weather and distance return expected wait time
in minutes
```

Use the UDF

```
describe ridetodest;
```

```
select user
, round(dist) as dist
, weather_description
, round(TAXI_WAIT(weather_description, dist)) as taxi_wait_min
from ridetodest emit changes;
```

```
select user
+ ' will be waiting ' + cast(round(TAXI_WAIT(weather_description, dist))
as varchar)
+ ' minutes for their trip of '
+ cast(round(dist) as varchar) +' km to ' + city_name
+ ' where it is ' + weather_description
from ridetodest emit changes;
```

```
Heidi will be waiting 14 minutes for their trip of 358 km to Bristol where
it is light rain
```

```
Bob will be waiting 4 minutes for their trip of 218 km to Manchester where
it is SUNNY
```

```
Frank will be waiting 15 minutes for their trip of 193 km to London where
it is heavy rain
```